

Case Studies in AI Uncertainty Quantification and Control

Edgar Dobriban


Department of Statistics and Data Science, the Wharton School, University of Pennsylvania

June 11, 2026



Context

- ▶ Artificial intelligence is emerging as an important technology
- ▶ Success stories: Large language models (ChatGPT, Claude), Image generation (Dall-E, Stable Diffusion, FLUX), Protein structure prediction ...



The screenshot shows a code editor with a file explorer on the left, a main code area, and a Claude Code sidebar on the right. The file explorer shows a project structure with folders like 'CP AUX LABELS' and 'code', and files like 'simulation.py'. The main code area shows the content of 'simulation.py', which is a Python script for a simulation. The Claude Code sidebar shows a search for 'class' definitions, a list of files found, and a summary of the codebase.

```
code > simulation.py > SimulationConfig
1 """
2 New simulation comparing three approaches with the same real calibration data:
3 1. Standard CP: 20 real calibration points
4 2. SPI: 20 real + 200 synthetic (K=10 per real point)
5 3. Flow-based CP: 20 real + 200 transformed synthetic + new synthetic for each predi
6 """
7 import numpy as np
8 import pandas as pd
9 import matplotlib.pyplot as plt
10 from dataclasses import dataclass
11 from typing import Dict
12 import time
13
14 # Import our modules
15 from data_generation import gen_real, gen_target, generate_synthetic_labels
16 from regression import fit_ols, score
```

Search(pattern: "class|def", path: "code")
Found 58 files (ctrl+r to expand)

Read(code/conformal_prediction.py)
Read 50 lines (ctrl+r to expand)


Read(experiments/README.md)
Read 223 lines (ctrl+r to expand)

This is a research codebase for conformal prediction with synthetic
implements and compares different methods for improving conformal pr
calibration data is limited but synthetic data is available.

Core Purpose

The codebase addresses the problem of small sample conformal predict
data. It compares 5 different approaches:

1. Standard CP: Baseline using only real data
2. Naive Pooling: Directly combining real + synthetic scores
3. ...



Solving a Research Problem in Mathematical Statistics with AI Assistance

Edgar Dobriban*

November 25, 2025

Abstract

Over the last few months, AI models including large language models have improved greatly. There are now several documented examples where they have helped professional mathematical scientists prove new results, sometimes even helping resolve known open problems. In this short note, we add another example to the list, by documenting how we were able to solve a previously unsolved research problem in robust mathematical statistics with crucial help from GPT-5. Our problem concerns robust density estimation, where the observations are perturbed by Wasserstein-bounded contaminations. In a previous preprint (Chao and Dobriban, 2023, arxiv:2308.01853v2), we have obtained upper and lower bounds on the minimax optimal estimation error; which were, however, not sharp.

Starting in October 2025, making significant use of GPT-5 Pro, we were able to derive the minimax optimal error rate (reported in version 3 of the above arxiv preprint). GPT-5 provided crucial help along the way, including by suggesting calculations that we did not think of, and techniques that were not familiar to us, such as the dynamic Benamou-Brenier formulation, for key steps in the analysis. Working with GPT-5 took a few weeks of effort, and we estimate that it could have taken several months to get the same results otherwise. At the same time, there are still areas where working with GPT-5 was challenging: it sometimes provided incorrect references, and glossed over details that sometimes took days of work to fill in. We outline our workflow and steps taken to mitigate issues. Overall, our work can serve as additional documentation for a new age of human-AI collaborative work in mathematical science.

AI as a Tool for Mathematics, Computer Science, and Machine Learning



Helping researchers integrate AI into core research workflows

📍 Seoul, South Korea

📅 July 10 or 11, 2026

Contact: ai4research-icml@googlegroups.com

Submit: [OpenReview](#)



Mikhail Belkin
UC San Diego



Sébastien Bubeck
OpenAI



Dmitriy Drusvyatskiy
UC San Diego



Ravi Vakil
Stanford University



Fanny Yang
ETH Zürich

STAI-X '26

Statistics and Trustworthy AI for Cross (X)-Domain Acceleration

To integrate statistics and artificial intelligence to accelerate trustworthy cross-domain discovery, innovation, and real-world impact.

Harvard University, Cambridge, MA • July 31 – August 1, 2026

Co-hosts

ASA StatsUpAI Interest Group (Lead Organizer) • Department of Statistics, Harvard University • Department of Biostatistics, Harvard University

Partner Organizations

National Academies Committee on Applied and Theoretical Statistics

Committee of Presidents of Statistical Societies

ASA



Partner Journals

JASA • AOAS • HDSR • ASA Discoveries • Canadian Journal of Statistics • Genome Research • GENETICS • Data Science in Science • Statistics and Data Science in Imaging



Xihong Lin
Harvard University



Wenyi Wang
MD Anderson Cancer Center



Linjun Zhang
Rutgers University





Tian Zheng
Columbia University

Current AI Has Problems


Science

Sycophantic AI decreases prosocial intentions and promotes dependence

MYRA CHENG , CINO O LEE , PRANAV KHADPE , SUNNY YU, DYLLAN HAN , AND DAN JURAFSKY  [Authors Info & Affiliations](#)



Recent Research



Statistical Methods in Generative AI

Edgar Dobriban¹


¹Department of Statistics and Data Science, University of Pennsylvania,
Philadelphia, PA, USA, 19104; email: dobriban@wharton.upenn.edu



[Home](#) / [A-Z Publications](#) / [Annual Review of Statistics and Its Application](#)

Annual Review of Statistics and Its Application

Uncertainty quantification and risk control




Imperfect Predictions

Predictors rarely achieve 100% accuracy



Image x



Predictor f

Incorrect!
(Ground truth label: $y = \text{"kite"}$)


"vulture"

Prediction
 $\hat{y} = \max_y f(y | x)$

Prediction Sets



Image x



Prediction Set $x \mapsto C(x)$

$\left\{ \begin{array}{l} \widehat{\text{vulture}}, \\ \text{kite}, \\ \text{bald eagle}, \\ \text{great grey owl}, \\ \text{bittern} \end{array} \right\}$


Output $C(x)$

Idea: Modify predictor f to predict **sets of labels**

Prediction Sets

Construct prediction sets based on a predictor f :

$$C_{\tau}(x) = \{y \mid f(y \mid x) \geq \tau\}$$



$$C_{\tau}(x) = \{\text{vulture, kite}\}$$


Case study I

Optimal decision-making based on prediction sets

Tao Wang and E.D.
ICML 2026 (oral presentation)

From prediction sets to decisions

- ▶ Prediction sets do not tell us what to do.
- ▶ A downstream decision has a **loss**: taking action a when the label is y costs $\ell(a, y)$.
- ▶ How to act when the label is Y plausibly in the prediction set?



Minimax policy

- ▶ For a prediction set $x \mapsto C(x)$, define the distributions compatible with conditional coverage

$$\mathcal{P}_\alpha(C) = \{P : P(Y \in C(X) \mid X = x) \geq 1 - \alpha \text{ for a.e. } x\}.$$

- ▶ **Minimax policy** π : minimizes risk over the worst-case distribution in this class

$$\inf_{\pi} \sup_{P \in \mathcal{P}_\alpha(C)} \mathbb{E}_P [\ell(\pi(C(X)), Y)].$$

Worst-case expected loss induced by a prediction set

For a fixed set $S \subseteq \mathcal{Y}$ and action a , write

$$\ell_S^{\text{in}}(a) = \sup_{y \in S} \ell(a, y), \quad \ell_S^{\text{out}}(a) = \sup_{y \notin S} \ell(a, y).$$

Worst-case risk for one set

$$L_S(a; \alpha) = \sup_{Q: Q(S) \geq 1-\alpha} \mathbb{E}_{Y \sim Q}[\ell(a, Y)] = \ell_S^{\text{in}}(a) + \alpha(\ell_S^{\text{out}}(a) - \ell_S^{\text{in}}(a))_+.$$

- ▶ Now consider any prediction set $x \mapsto C(x)$
- ▶ We then show that the minimax policy reduces to a pointwise decision:


$$\pi^*(C(x)) \in \arg \min_{a \in \mathcal{A}} L_{C(x)}(a; \alpha).$$

- ▶ Can use this policy for any prediction set C

ROCP: conformalized risk-optimal sets

- ▶ We then characterize oracle optimal prediction sets, following Kiyani et al., 2025.
- ▶ **Empirically:** Estimate the components of optimal prediction sets based on the loss and classifier.
- ▶ Conformalize them to obtain C_{ROCP} and act with $\arg \min_a L_{C_{\text{ROCP}}(X_{\text{test}})}(a; \alpha)$.


Medical diagnosis experiment: setup



Data and model

- COVID-19 Radiography Database.


Medical diagnosis experiment: setup



Data and model

- ▶ COVID-19 Radiography Database.
- ▶ Labels: $\mathcal{Y} = \{\text{Normal, Pneumon., COVID, Opacity}\}.$

Medical diagnosis experiment: setup



Data and model

- ▶ COVID-19 Radiography Database.
- ▶ Labels: $\mathcal{Y} = \{\text{Normal, Pneumon., COVID, Opacity}\}$.
- ▶ Predictor: Inception-v3 fine-tuned on radiographs.

Medical diagnosis experiment: setup

COVID-19



Viral Pneumonia



Normal



Actions and losses


Actions are

$$\mathcal{A} = \{\text{NoAct}, \text{Antibiotics}, \text{Quarantine}, \text{Test}\}.$$

Data and model

- ▶ COVID-19 Radiography Database.
- ▶ Labels: $\mathcal{Y} = \{\text{Normal}, \text{Pneumon.}, \text{COVID}, \text{Opacity}\}.$
- ▶ Predictor: Inception-v3 fine-tuned on radiographs.

Medical diagnosis experiment: setup



Actions and losses

Actions are

$$\mathcal{A} = \{\text{NoAct}, \text{Antibiotics}, \text{Quarantine}, \text{Test}\}.$$


Loss matrix:

	NoAct	Antib.	Quar.	Test
Normal	0	8	8	6
Pneum.	100	0	70	3
COVID	100	70	0	2
Opacity	90	60	60	0

Data and model

- ▶ COVID-19 Radiography Database.
- ▶ Labels: $\mathcal{Y} = \{\text{Normal}, \text{Pneumon.}, \text{COVID}, \text{Opacity}\}.$
- ▶ Predictor: Inception-v3 fine-tuned on radiographs.

Medical diagnosis experiment: setup



Actions and losses

Actions are

$$\mathcal{A} = \{\text{NoAct}, \text{Antibiotics}, \text{Quarantine}, \text{Test}\}.$$

Loss matrix:

	NoAct	Antib.	Quar.	Test
Normal	0	8	8	6
Pneum.	100	0	70	3
COVID	100	70	0	2
Opacity	90	60	60	0

Data and model

- ▶ COVID-19 Radiography Database.
- ▶ Labels: $\mathcal{Y} = \{\text{Normal}, \text{Pneumon.}, \text{COVID}, \text{Opacity}\}.$
- ▶ Predictor: Inception-v3 fine-tuned on radiographs.

ROCP is compared to

- ▶ RAC (Kiyani et al., 2025),
- ▶ LAS: Robust or max-min actions using CP with score $[f(x)]_{\mathcal{Y}}.$
- ▶ Best-response: $\arg \min \mathbb{E}_{Y \sim f(x)} \ell(a, Y)$

Medical diagnosis experiment: results




Figure: (a) Worst-case loss;

- In contrast to RAC, ROCP stays stable because it accounts for the α fraction of out-of-set mass.

Medical diagnosis experiment: results




Figure: (a) Worst-case loss; (b) critical mistake rates.

- ▶ In contrast to RAC, ROCP stays stable because it accounts for the α fraction of out-of-set mass.
- ▶ ROCP reduces worst-case loss, realized loss, and critical mistakes.

Case study II

Online conformal prediction via universal portfolio algorithms

Tuo Liu, E.D, and Francesco Orabona
ICML 2026 (Spotlight)

Online conformal prediction: setup

- ▶ At time t , observe features x_t . Predict outcome Y_t .
- ▶ Construct prediction set

$$\widehat{C}_t = [\widehat{Y}_t - b_t, \widehat{Y}_t + b_t].$$

- ▶ After observing Y_t , define the nonconformity score $S_t = |Y_t - \widehat{Y}_t|$.


Goal

For arbitrary, possibly adversarial streams, want the time-averaged fraction of coverage events $Y_t \in \widehat{C}_t$ to approach $1 - \alpha$. Equivalently,

$$\frac{1}{T} \sum_{t=1}^T \mathbf{1}\{S_t \leq b_t\} \rightarrow 1 - \alpha.$$

This can be viewed as online learning of the $(1 - \alpha)$ -quantile of the changing score sequence $S_1, S_2, \dots$


— Nonconformity Scores



Review: Pinball/quantile loss

Quantile loss

$$\ell_t(b) = \ell^{(1-\alpha)}(b, S_t) = (1-\alpha)(S_t - b)_+ + \alpha(b - S_t)_+.$$



Use the subgradient

$$g_t \in \partial \ell_t(b_t), \quad g_t = \begin{cases} -(1-\alpha), & S_t > b_t, \quad (\text{miscoverage}), \\ \alpha, & S_t \leq b_t, \quad (\text{coverage}). \end{cases}$$

Then the miscoverage is (Gibbs & Candès, '21, Angelopoulos, Jordan & Tibshirani, '25, etc)

$$\frac{1}{T} \sum_{t=1}^T \mathbf{1}\{S_t \leq b_t\} - (1-\alpha) = \frac{1}{T} \sum_{t=1}^T g_t.$$

Classical and linear regret

- ▶ Classical regret at comparator u :

$$\text{Regret}_T(u) := \sum_{t=1}^T \ell_t(b_t) - \sum_{t=1}^T \ell_t(u) .$$

- ▶ **Linearized regret** (Gordon '99, Zinkevich '03, etc) at comparator u (for $g_t \in \partial \ell_t(b_t)$) sums the linearizations $g_t(b_t - u)$ of the loss differences $\ell_t(b_t) - \ell_t(u)$ around b_t :

$$\text{LinRegret}_T(u) = \sum_{t=1}^T g_t(b_t - u).$$

- ▶ By convexity, $\text{Regret}_T(u) \leq \text{LinRegret}_T(u)$ for all u .

A black-box regret-to-coverage theorem

From linearized regret to coverage (Liu, D., and Orabona, '26)

For any upper bound F_T on LinRegret,

$$-\sum_{t=1}^T g_t \in \{z : F_T^*(z) \leq (1 - \alpha) \sum_{t=1}^T S_t\},$$

where F_T^* is the Fenchel conjugate of F_T : $F_T^*(z) = \sup_u \{z^\top u - F_T(u)\}$.


- ▶ Any algorithm with a bounded linearized regret inherits a coverage bound.
- ▶ Ordinary regret is insufficient (Angelopoulos et al., '25).

Conformal prediction via universal portfolio algorithms (UP-OCP)

- ▶ Next: choose methods for online learning with a good control of linearized regret.
- ▶ **Equivalence of coin betting & Online Conv Optimization**: it is enough to achieve high wealth while betting in the "conformal market" with two synthetic stocks and returns

$$w_{t,1} = 1 - \frac{g_t}{\alpha}, \quad w_{t,2} = 1 + \frac{g_t}{1 - \alpha}.$$

- Stock 1 pays off on miscoverage rounds.
- Stock 2 pays off on coverage rounds.
- ▶ **Universal portfolio algorithms** (Cover and Ordentlich, 2002) have near-optimal wealth.



UP-OCP: full online algorithm

State and prediction

Input: target miscoverage rate $\alpha \in (0, 1)$.

Initialize: wealth $W_0 \leftarrow 1$, miscoverage count $N \leftarrow 0$.

1. For $t = 1, 2, \dots$, set

$$\lambda_t \leftarrow \frac{N + 1/2}{t}, b_t \leftarrow \max \left\{ 0, W_{t-1} \frac{\lambda_t - \alpha}{\alpha(1 - \alpha)} \right\}.$$

2. Output the interval $\hat{C}_t \leftarrow [\hat{Y}_t - b_t, \hat{Y}_t + b_t]$.

3. Observe Y_t and compute $S_t \leftarrow |Y_t - \hat{Y}_t|$.

Multiplicative update

$$\text{miscov. : } N \leftarrow N + 1, W_t \leftarrow W_{t-1} \frac{\lambda_t}{\alpha},$$

$$\text{coverage : } W_t \leftarrow W_{t-1} \frac{1 - \lambda_t}{1 - \alpha}.$$

- ▶ $\lambda_t \in [0, 1]$ is the universal portfolio weight (alloc. to "stock" 1).
- ▶ Requires **no hyperparameter tuning**.

Finite-time coverage for UP-OCP

Assume at most polynomial score growth, $S_t \leq Dt^q$. Define

$$\varepsilon_T = \frac{1}{T} \left[\log \left(1 + \frac{(1-\alpha)D(\textcolor{red}{T} + 1)^{q+1}}{q+1} \right) + \frac{1}{2} \log(\pi(T+1)) \right].$$

Bound on mis-coverage (Liu, D., and Orabona, '26)

The miscoverage satisfies

$$\text{MisCov}_T \leq \varepsilon_T + \sqrt{2\alpha(1-\alpha)\varepsilon_T} = O \left(\frac{\log T}{T} + \sqrt{\alpha \frac{\log T}{T}} \right).$$

The variance-adaptive term improves as α approaches 0.

Empirical experiment

Data and base forecasts

- ▶ Daily log American Express stock prices;
- ▶ Prophet (dev. by Meta) re-fit daily.

Empirical experiment

Data and base forecasts

- ▶ Daily log American Express stock prices;
- ▶ Prophet (dev. by Meta) re-fit daily.

Metrics and baselines

- ▶ Show tradeoff of realized coverage versus average interval size.

Empirical experiment

Data and base forecasts

- ▶ Daily log American Express stock prices;
- ▶ Prophet (dev. by Meta) re-fit daily.

Metrics and baselines

- ▶ Show tradeoff of realized coverage versus average interval size.
- ▶ Baselines:
 - KT (Podkopaev et al., 2024)
 - DtACI (Gibbs and Candès [2024])
 - scale-free OGD (Orabona and Pál, 2018)
 - conformal P/PI control (Angelopoulos et al., 2023).

Empirical experiment



Data and base forecasts

- ▶ Daily log American Express stock prices;
- ▶ Prophet (dev. by Meta) re-fit daily.

Metrics and baselines

- ▶ Show tradeoff of realized coverage versus average interval size.
- ▶ Baselines:
 - KT (Podkopaev et al., 2024)
 - DtACI (Gibbs and Candès [2024])
 - scale-free OGD (Orabona and Pál, 2018)
 - conformal P/PI control (Angelopoulos et al., 2023).
- ▶ Tuned baselines get ex-post grid search; UP-OCF is parameter-free.

Empirical results



► UP-0CP gives the best observed size/coverage tradeoff.

Case study III

Risk-controlled post-processing of decision policies

Sunay Joshi, Tao Wang, Hamed Hassani, and E.D.

How to minimally post-process a policy to ensure risk control?

- ▶ Consider a setting where there is a **baseline policy** (LLM, patient vaccination procedure, etc).
- ▶ However, new developments require us to satisfy a **new constraint** (e.g., cost, safety).
- ▶ We want to stick to our existing policy as much as possible, while satisfying the constraint.
- ▶ The constraint is on some downstream loss.

Mathematical setup

Given a deterministic baseline policy π_0 , loss ℓ , and cutoff c , define the violation event

$$L_\pi(X, Y) = \mathbf{1}\{\ell(\pi(X), Y) \geq c\}.$$

How to minimally post-process a policy to ensure risk control at budget ε ? Solve:

$$\max_{\pi} \mathbb{P}\{\pi(X) = \pi_0(X)\} \quad \text{s.t.} \quad \mathbb{E}L_\pi(X, Y) \leq \varepsilon.$$

Population oracle: switch on high-risk-reduction contexts

Let

$$g(a, x) = \mathbb{P}\{\ell(a, Y) \geq c \mid X = x\}, \quad \pi_*(x) \in \arg \min_a g(a, x),$$

and define the violation-risk reduction score


$$\Delta(x) = g(\pi_0(x), x) - g(\pi_*(x), x).$$

Threshold structure


The optimal policy follows π_0 on low-score contexts and switches to π_* on high-score contexts:

$$\pi^*(x; \tau) = \begin{cases} \pi_0(x), & \Delta(x) < \tau, \\ \pi_*(x), & \Delta(x) \geq \tau. \end{cases}$$

The threshold is chosen so the risk constraint is tight whenever possible.



Finite-sample post-processing algorithm



The fitted switching policy is

$$\hat{\pi}(x; \tau) = \begin{cases} \pi_0(x), & \hat{\Delta}(x) < \tau, \\ \hat{\pi}_*(x), & \hat{\Delta}(x) \geq \tau. \end{cases}$$

For calibration data $Z_i = (X_i, Y_i)$ and grid $\mathcal{T}_n = \{0, \top\} \cup \{\hat{\Delta}(X_i) : i \in [n]\}$, define

$$\hat{R}_n^+(\tau) = \frac{1}{n+1} \left(\sum_{i=1}^n \mathbf{1}\{\ell(\hat{\pi}(X_i; \tau), Y_i) \geq c\} + 1 \right).$$

Choose

$$\hat{\tau} = \max\{\tau \in \mathcal{T}_n : \hat{R}_n^+(\tau) \leq \varepsilon\}.$$

- ▶ The largest feasible threshold preserves as much agreement with π_0 as possible.
- ▶ The +1 term is the conformal-risk-control finite-sample correction.

Guarantees: exact-safe and general fitted fallbacks

Exact-safe fallback

If the fallback is exact-safe, $\ell(\pi_{\text{safe}}(x), y) < c$ for all x, y , then under exchangeability

$$\mathbb{E}L(Z_{n+1}; \hat{\tau}) \leq \varepsilon.$$

This is finite-sample and distribution-free.


Fitted fallback

For a general fitted fallback, the loss can be non-monotone in τ . Rank-stability analysis gives

$$\mathbb{E}L(Z_{n+1}; \hat{\tau}) \leq \varepsilon + O\left(\frac{\log n}{n}\right).$$

When a exact-safe fallback is available, the conformalized policy preserves **nearly as much agreement with π_0** as the population oracle, up to finite-sample and score-estimation terms.

Synthetic experiment: targeted switching



- The synthetic task isolates the oracle threshold structure.
- The learned threshold intervenes selectively, not by score-blind random mixing.
- Risk budgets are met or nearly met while preserving substantially more agreement with the baseline.

Empirical examples: clinical and LLM routing

COVID radiograph diagnosis

- ▶ Same four labels/actions as the ROCP medical example.
- ▶ Split: 70% training, 10% baseline calibration, 10% threshold calibration, 10% test.
- ▶ If $c > 6$, *Additional Testing* is exact-safe; at $c = 7, \varepsilon = 0.02$, switching only 6% of cases attains violation risk 0.02.
- ▶ For $c \leq 6$, no exact-safe action exists; the fitted fallback regime still tracks the target risk up to stability slack.

LLM thinking-mode routing

- ▶ Baseline: Qwen3-4B or Qwen3-1.7B in non-thinking mode.
- ▶ Fallback: Qwen3-32B in thinking mode on MMLU-Pro.
- ▶ Constraint: answer-error rate $\mathbb{P}(\pi(X) \neq Y) \leq \varepsilon$.
- ▶ At matched error risk, targeted routing switches less often than random mixing and saves 12–19% FLOPs.

Case study IV

MultiRisk: multiple risk control via iterative score thresholding

Sunay Joshi, Yan Sun, Hamed Hassani, and E.D.

Multiple risks at test time

For an input X , a model first generates an output Y . We then evaluate prioritized scores $S_1, \dots, S_m$ and compare them with thresholds $\lambda_1, \dots, \lambda_m$.

Priority rule

If every score passes, return the original output Y . Otherwise let

$$j^* = \min\{j : S_j(X, Y, Y^*) > \lambda_j\}$$

and return the behavior $B_{j^*}(X, Y)$.


- ▶ Behaviors can be refusal, resampling, warning, routing, or ordinary pass-through.
- ▶ The priority order makes the outcome unambiguous: exactly one behavior occurs.
- ▶ Larger thresholds allow more responses through; smaller thresholds intervene earlier.

Example control flow: LLM abstention

Threshold selection



Sequential routing



Sequential multiple risk control problem

Constraint loss for behavior j is the cost of activating B_j after all earlier filters pass:

$$L_j(\lambda_{1:j}) = V_j(X, Y, Y^*) \mathbf{1}\{S_1 \leq \lambda_1, \dots, S_{j-1} \leq \lambda_{j-1}, S_j > \lambda_j\}.$$

Objective loss is the cost of returning the original output if all filters pass:

$$L_{m+1}(\lambda_{1:m}) = V_{m+1}(X, Y, Y^*) \mathbf{1}\{S_1 \leq \lambda_1, \dots, S_m \leq \lambda_m\}.$$

Population problem

$$\min_{\lambda_1 \in \Lambda_1, \dots, \lambda_m \in \Lambda_m} \mathbb{E} L_{m+1}(\lambda_{1:m}) \quad \text{s.t.} \quad \mathbb{E} L_j(\lambda_{1:j}) \leq \beta_j \quad \forall j \in [m].$$

- ▶ Each β_j is a provider-chosen risk budget for one intervention.
- ▶ The objective is expected cost of ordinary pass-through.
- ▶ The risks interact because later filters see only the examples that earlier filters did not catch.


Algorithmic idea: dynamic programming with guardrailed thresholds

Idea: solve empirical constraint risk equations $\widehat{R}_j(\lambda_j; \lambda_{1:(j-1)}) = \beta_j$ **sequentially**, for $j = 1, 2, \dots, m$.

Conformal correction: given earlier thresholds $\lambda_{1:(j-1)}$, MultiRisk uses a **conservative empirical risk function**

$$\widehat{R}_j^+(\lambda_j; \lambda_{1:(j-1)}) = \frac{1}{n+1} \left\{ \sum_{i=1}^n L_j^{(i)}(\lambda_{1:j}) + L_j^{\max} \right\}.$$

- ▶ Here $L_j^{\max}$ is an a.s. upper bound on L_j .
- ▶ MultiRisk-Base solves the empirical dynamic program using inverse risk curves.
- ▶ MultiRisk adds the $L_j^{\max}/(n+1)$ bump needed for finite-sample control.
- ▶ Auxiliary upper thresholds and budget shrinkage keep earlier decisions valid when later thresholds are chosen.



Recursive inequalities used to lower bound risk tightness

calibration scores and costs $\rightarrow \widehat{\lambda}_1, \dots, \widehat{\lambda}_m$

MultiRisk guarantees

Risk control

Under exchangeability and bounded costs,

$$\mathbb{E}L_j^{(n+1)}(\hat{\lambda}_{1:j}) \leq \beta_j$$

for all $j \in [m]$.

Near-tightness

Under feasibility and continuity, for some constant A_j ,

$$\mathbb{E}L_j^{(n+1)}(\hat{\lambda}_{1:j}) \geq \beta_j - \frac{A_j}{n+1}$$

for all $j \in [m]$.

Near-optimal objective

Under empirical-process regularity, the achieved objective nearly matches the oracle.

- ▶ Risk control is **distribution-free**: the calibration and test points only need exchangeability.
- ▶ Tightness says the budgets are nearly exhausted.
- ▶ Near-optimality explains why MultiRisk can outperform conservative multiple-testing baselines.

LLM alignment experiment: setup

Data and model


- ▶ Dataset: PKU-SafeRLHF-30K, with calibration and test sets of size 500.
- ▶ Generator: ALPACA-7B-REPRODUCED with greedy decoding.
- ▶ If any threshold is exceeded, the deployed behavior is abstention.

Scores and costs


- ▶ S_1 : Llama-Guard probability for violent crimes, weapons, or self-harm.
- ▶ S_2 : Llama-Guard probability for non-violent crimes, defamation, or privacy.
- ▶ S_3 : perplexity of the generated response; objective cost is shifted unhelpfulness.

Choose three thresholds to **minimize unhelpfulness** while controlling the **abstention-triggering risks**.

Empirical illustration: LLM alignment with three constraints



objective surface



- ▶ MultiRisk controls the three risks close to their target budgets across the grid.
- ▶ Joint thresholding improves the objective relative to Learn-then-Test with Bonferroni correction.
- ▶ MultiRisk-Base is useful empirically, while MultiRisk has finite-sample guarantees.

Summary

- ▶ We discussed several case studies in AI uncertainty quantification and control: starting from AI outputs and post-processing to ensure guarantees such as coverage or risk control.
- ▶ Thank you!

References

- ▶ Tao Wang and E.D. *Optimal Decision-Making Based on Prediction Sets*. ICML 2026.
- ▶ Tuo Liu, E.D, and Francesco Orabona. *Online Conformal Prediction via Universal Portfolio Algorithms*. ICML 2026.
- ▶ Sunay Joshi, Tao Wang, Hamed Hassani, and E.D. *Risk-Controlled Post-Processing of Decision Policies*. arXiv:2605.06479.
- ▶ Sunay Joshi, Yan Sun, Hamed Hassani, and E.D. *MultiRisk: Multiple Risk Control via Iterative Score Thresholding*. arXiv:2512.24587.

