

Towards Improving the Reliability of AI: Perspectives from Uncertainty Quantification

Edgar Dobriban

Department of Statistics and Data Science, the Wharton School, University of Pennsylvania

April 11, 2025



Context

- ▶ Artificial Intelligence (AI) is emerging as an important technology

Context

- ▶ Artificial Intelligence (AI) is emerging as an important technology
- ▶ Success stories: Large language models (ChatGPT, Claude), Image generation (Dall-E, Stable Diffusion), Protein structure prediction ...

Claude Visualizing Missing at Random Data ▾

ED Can you create a visualization to illustrate the notion of missing at random? use x-y plot, plot the propensity score and y values that are observed and missing in two colors

Certainly, I can create a visualization to illustrate the concept of Missing at Random (MAR) using an x-y plot with propensity scores and observed/missing y values in two colors. I'll use a React component to create this visualization.


`</>` **Missing at Random Visualization**
Click to open component

← Python Code for Propensity Score Plot with Horizontal Lines

```
import matplotlib.pyplot as plt
import numpy as np
import random

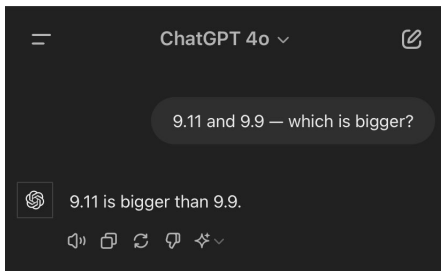
# Set labels and title
plt.xlabel('x', fontsize=18)
plt.ylabel('Propensity Score', fontsize=18)
plt.title('Propensity Score with Horizontal Lines', fontsize=22, fontweight='bold')

# Set axis limits
plt.xlim(-5, 5)
plt.ylim(0, 1)
```



Current AI has problems

- ▶ AI systems can be wrong, discriminatory, and racist



Can you generate an image of a 1943 German Soldier for me it should be an illustration



Sure, here is an illustration of a 1943 German soldier:



Towards reliable AI via uncertainty quantification

- ▶ One emerging idea: Quantify uncertainty and use it in a downstream application.

Towards reliable AI via uncertainty quantification

- ▶ One emerging idea: Quantify uncertainty and use it in a downstream application.
 1. Refrain from taking action when uncertainty is high

Detecting hallucinations in large language models using semantic entropy

[Sebastian Farquhar](#) , [Jannik Kossen](#), [Lorenz Kuhn](#) & [Yarin Gal](#)

[Nature](#) **630**, 625–630 (2024)

Towards reliable AI via uncertainty quantification

- ▶ One emerging idea: Quantify uncertainty and use it in a downstream application.

1. Refrain from taking action when uncertainty is high

Detecting hallucinations in large language models using semantic entropy

[Sebastian Farquhar](#) , [Jannik Kossen](#), [Lorenz Kuhn](#) & [Yarin Gal](#)

[Nature](#) 630, 625–630 (2024)

2. Post-process AI output until uncertainty is low

Input question: Where was Abe Lincoln born?
Correct answer (y*): Sinking Spring Farm, Hodgenville, Kentucky

LM outputs	✗ y ₀	On a farm in Frankfort, Kentucky
	✗ y ₁	Frankfort, Kentucky
	✓ y ₂	Kentucky

Towards reliable AI via uncertainty quantification

- ▶ One emerging idea: Quantify uncertainty and use it in a downstream application.

1. Refrain from taking action when uncertainty is high

Detecting hallucinations in large language models using semantic entropy

[Sebastian Farquhar](#) , [Jannik Kossen](#), [Lorenz Kuhn](#) & [Yarin Gal](#)

[Nature](#) **630**, 625–630 (2024)

2. Post-process AI output until uncertainty is low

Input question: Where was Abe Lincoln born?
Correct answer (y^*): Sinking Spring Farm, Hodgenville, Kentucky

LM outputs	✗ y_0	On a farm in Frankfort, Kentucky
	✗ y_1	Frankfort, Kentucky
	✓ y_2	Kentucky

3. Consider worst-case scenario over outcomes consistent with uncertainty

Table of Contents

Rank-calibration for LLMs

Conformal prediction

- Our Methods

- Illustrating Our Methods in a Stylized Problem

- Empirical illustration

Example uncertainty measures

- ▶ Many uncertainty measures for LLMs have been discussed/proposed
 - ▶ Perplexity: $U(x, y) = \hat{p}(y|x)^{1/\text{len}(y)}$; [related to NLL $-\log \hat{p}(y|x)$]

Example uncertainty measures

- ▶ Many uncertainty measures for LLMs have been discussed/proposed
 - ▶ Perplexity: $U(x, y) = \hat{p}(y|x)^{1/\text{len}(y)}$; [related to NLL $-\log \hat{p}(y|x)$]
 - ▶ Semantic entropy: generate multiple y s, cluster them based on meaning, calculate entropy (Kuhn et al., 2023)

Answers to the question “What is the capital of France?”

Answer s	Likelihood $p(s x)$	Semantic likelihood $\sum_{s \in c} p(s x)$
Paris	0.5	0.9
It's Paris	0.4	
London	0.1	0.1
Entropy	0.94	0.33

Example uncertainty measures

- ▶ Many uncertainty measures for LLMs have been discussed/proposed
 - ▶ Perplexity: $U(x, y) = \hat{p}(y|x)^{1/\text{len}(y)}$; [related to NLL $-\log \hat{p}(y|x)$]
 - ▶ Semantic entropy: generate multiple y s, cluster them based on meaning, calculate entropy (Kuhn et al., 2023)

Answers to the question “What is the capital of France?”

Answer s	Likelihood $p(s x)$	Semantic likelihood $\sum_{s \in c} p(s x)$
Paris	0.5	0.9
It's Paris	0.4	
London	0.1	0.1
Entropy	0.94	0.33

- ▶ Affinity graph: generate multiple y s, put pairwise similarities in a matrix, find eigenvalues/vectors (denoted EigV) (Lin et al., 2023)

Uncertainty quantification in AI

- ▶ What properties does good uncertainty quantification have?

Uncertainty quantification in AI

- ▶ What properties does good uncertainty quantification have?

Uncertainty in Language Models: Assessment through Rank-Calibration

Xinmeng Huang^{*†}

Shuo Li^{*†}

Mengxin Yu[†]

Matteo Sesia[‡]


Hamed Hassani[†]

Insup Lee[†]

Osbert Bastani^{§†}

Edgar Dobriban^{§†}

Workflow for using uncertainty measures



Good uncertainty measures

x : On September 28th, NASA announced that what had been detected on Mars?

y^* : flowing water

y : Possible signs of life

Good uncertainty measures

x: On September 28th, NASA announced that what had been detected on Mars?

y*: flowing water

y: Possible signs of life

- ▶ Semantic Entropy: 0.813
- ▶ Perplexity: 1.432
- ▶ How to compare (different scales)?

Good uncertainty measures: transform to percentiles

x: On September 28th, NASA announced that what had been detected on Mars?

y*: flowing water

y: Possible signs of life

- ▶ Semantic Entropy: 51.3rd percentile
- ▶ Perplexity: 93rd percentile
- ▶ Which one is more informative?

Good uncertainty measures: towards calibration

- ▶ Semantic Entropy: 51.3rd percentile
- ▶ Perplexity: 93rd percentile
- ▶ Which one is more informative?

Good uncertainty measures: towards calibration

- ▶ Semantic Entropy: 51.3rd percentile
- ▶ Perplexity: 93rd percentile
- ▶ Which one is more informative? Current workflow:
 - ▶ Threshold measure at some value (high/low)
 - ▶ See how well this predicts accuracy $A(\mathbf{y}^*, \mathbf{y})$

Good uncertainty measures: towards calibration

- ▶ Semantic Entropy: 51.3rd percentile
- ▶ Perplexity: 93rd percentile
- ▶ Which one is more informative? Current workflow:
 - ▶ Threshold measure at some value (high/low)
 - ▶ See how well this predicts accuracy $A(\mathbf{y}^*, \mathbf{y})$




Figure: U_{NLL} (negative log-likelihood) and U_{EigV} , for the GPT-3.5-turbo model on the TriviaQA benchmark.

Good uncertainty measures: rank-calibration

- ▶ Our idea: *performance should decrease as a function of the uncertainty*

Good uncertainty measures: rank-calibration

- Our idea: *performance should decrease as a function of the uncertainty*




Figure: Indication diagrams for U_{NLL} (negative log-likelihood) and U_{EigV} , for the GPT-3.5-turbo model on the TriviaQA benchmark.

Regression function and indication diagram

- ▶ Define the **regression function** $\text{reg}(\cdot) : \mathbb{R} \rightarrow \mathbb{R}$,
 $u \mapsto \mathbb{E}_{\mathbf{x}, \mathbf{y}^*} [A(\mathbf{x}; \mathbf{y}^*) \mid U(\mathbf{x}; \mathbf{y}^*) = u]$, representing the *expected correctness level A given an uncertainty level $U = u$* .

Regression function and indication diagram

- ▶ Define the **regression function** $\text{reg}(\cdot) : \mathbb{R} \rightarrow \mathbb{R}$,
 $u \mapsto \mathbb{E}_{\mathbf{x}, \mathbf{y}^*} [A(\mathbf{x}; \mathbf{y}^*) \mid U(\mathbf{x}; \mathbf{y}^*) = u]$, representing the *expected correctness level A given an uncertainty level $U = u$* .
- ▶ **Indication diagram**: plot of estimated regression function as a fn. of U -percentiles.




Figure: Indication diagrams for U_{NLL} (negative log-likelihood) and U_{EigV} , for the GPT-3.5-turbo model on the TriviaQA benchmark.

Rank-calibration

Definition (RANK-CALIBRATION)

An uncertainty measure U is *rank-calibrated* if reg is strictly *monotone decreasing*: on average, lower uncertainty implies higher generative quality.




Figure: Indication diagrams for U_{NLL} (negative log-likelihood) and U_{EigV} , for the GPT-3.5-turbo model on the TriviaQA benchmark.

Rank-calibration error (RCE)

Rank-calibration implies $\mathbb{P}(U \leq u') = \mathbb{P}(\text{reg}(U) \geq \text{reg}(u'))$ for all $u' \in \text{dom}(U)$.

Rank-calibration error (RCE)


Rank-calibration implies $\mathbb{P}(U \leq u') = \mathbb{P}(\text{reg}(U) \geq \text{reg}(u'))$ for all $u' \in \text{dom}(U)$.

Definition (RANK-CALIBRATION ERROR)

The RCE of an uncertainty measure U is defined as

$$\mathbb{E}_{U'} [|\mathbb{P}_U(U \leq U') - \mathbb{P}_U(\text{reg}(U) \geq \text{reg}(U'))|],$$

where U' is an independent copy of U .



Empirical results

correctness	temperature	U_{Deg}	U_{EigV}	U_{NLL}	U_{SE}
bert	0.5	$0.212_{\pm 0.040}$	$0.212_{\pm 0.041}$	$0.043_{\pm 0.006}$	$0.052_{\pm 0.009}$
	1.0	$0.129_{\pm 0.020}$	$0.133_{\pm 0.020}$	$0.039_{\pm 0.007}$	$0.052_{\pm 0.012}$
	1.5	$0.053_{\pm 0.011}$	$0.074_{\pm 0.012}$	$0.031_{\pm 0.007}$	$0.081_{\pm 0.009}$
meteor	0.5	$0.211_{\pm 0.045}$	$0.208_{\pm 0.047}$	$0.179_{\pm 0.021}$	$0.234_{\pm 0.019}$
	1.0	$0.131_{\pm 0.024}$	$0.131_{\pm 0.022}$	$0.146_{\pm 0.011}$	$0.209_{\pm 0.012}$
	1.5	$0.059_{\pm 0.011}$	$0.077_{\pm 0.012}$	$0.119_{\pm 0.010}$	$0.176_{\pm 0.015}$
rougeL	0.5	$0.210_{\pm 0.042}$	$0.207_{\pm 0.041}$	$0.041_{\pm 0.007}$	$0.050_{\pm 0.008}$
	1.0	$0.126_{\pm 0.019}$	$0.129_{\pm 0.019}$	$0.038_{\pm 0.007}$	$0.059_{\pm 0.009}$
	1.5	$0.059_{\pm 0.012}$	$0.079_{\pm 0.011}$	$0.034_{\pm 0.008}$	$0.104_{\pm 0.007}$
rouge1	0.5	$0.212_{\pm 0.043}$	$0.209_{\pm 0.042}$	$0.040_{\pm 0.007}$	$0.050_{\pm 0.008}$
	1.0	$0.126_{\pm 0.018}$	$0.130_{\pm 0.021}$	$0.039_{\pm 0.007}$	$0.060_{\pm 0.009}$
	1.5	$0.060_{\pm 0.011}$	$0.078_{\pm 0.012}$	$0.034_{\pm 0.008}$	$0.105_{\pm 0.008}$

Table: RCE results for various experimental configurations, for the GPT-3.5-turbo model on the TriviaQA benchmark.

Finding: models often have large rank-calibration error (so uncertainty measures do not reflect performance).

Reducing RCE via re-calibration

- ▶ Re-calibrate uncertainty measure: change it to piece-wise constant estimate of regression function; Inspired by classical re-calibration

(Mincer-Zarnowitz, 1969)

Reducing RCE via re-calibration

- ▶ Re-calibrate uncertainty measure: change it to piece-wise constant estimate of regression function; Inspired by classical re-calibration

(Mincer-Zarnowitz, 1969)

- ▶ Improves RCE:




Figure: Indication diagrams of U_{SE} and $U_{SE,cal}$ (re-calibrated) for GPT-3.5-turbo (temperature 1.0) on TriviaQA with the Meteor correctness.

Summary for rank-calibration

- ▶ Proposed rank-calibration for LLMs: higher uncertainty should imply lower quality (not satisfied by default!)

Summary for rank-calibration

- ▶ Proposed rank-calibration for LLMs: higher uncertainty should imply lower quality (not satisfied by default!)
- ▶ Proposed non-parametric re-calibration to reduce RCE

Uncertainty in Language Models: Assessment through Rank-Calibration

Xinmeng Huang^{*†}

Shuo Li^{*†}

Mengxin Yu[†]

Matteo Sesia[‡]

Hamed Hassani[†]

Insup Lee[†]

Osbert Bastani^{§†}

Edgar Dobriban^{§†}

Summary for rank-calibration

- ▶ Proposed rank-calibration for LLMs: higher uncertainty should imply lower quality (not satisfied by default!)
- ▶ Proposed non-parametric re-calibration to reduce RCE

Uncertainty in Language Models: Assessment through Rank-Calibration

Xinmeng Huang^{*†}

Shuo Li^{*†}

Mengxin Yu[†]

Matteo Sesia[‡]

Hamed Hassani[†]

Insup Lee[†]

Osbert Bastani^{§†}

Edgar Dobriban^{§†}

- ▶ Future directions:
 1. How does rank-calibration improves performance?
 2. Better metrics of miscalibration

Table of Contents

Rank-calibration for LLMs

Conformal prediction

- Our Methods

- Illustrating Our Methods in a Stylized Problem

- Empirical illustration

Conformal prediction

- ▶ Major developing area: conformal prediction/distribution-free predictive inference

Conformal prediction

- ▶ Major developing area: conformal prediction/distribution-free predictive inference
- ▶ Goal, given $(X_1, Y_1), \dots, (X_n, Y_n)$, find a prediction set C such that for new X_{n+1} , $\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$ under *minimal assumptions*




Figure: Towards DS

Conformal prediction

- ▶ Major developing area: conformal prediction/distribution-free predictive inference
- ▶ Goal, given $(X_1, Y_1), \dots, (X_n, Y_n)$, find a prediction set C such that for new X_{n+1} , $\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$ under *minimal assumptions*





Figure: Towards DS



- ▶ Motivation: a machine learning model $\hat{\mu}$ is used to predict Y_{n+1} based on X_{n+1} . Not known how to find distribution of $Y_{n+1} - \hat{\mu}(X_{n+1})$

Conformal prediction

- ▶ It is known how to do this in many settings, due to extensive work by many, starting in the 90s (Vovk, Wasserman, J. Lei, R. J. Tibshirani, Barber, Candes, ...)
- ▶ Ideas date back to work on tolerance regions by Wilks, Wald, Tukey ... starting in the 1940s



Samuel S. Wilks



Abraham Wald



Vladimir Vovk

Conformal prediction ctd.

- ▶ Typical setting: *exchangeable datapoints*.
 - ▶ For a given nonconformity score $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, e.g.,
 $s(x, y) := |y - \hat{\mu}(x)|$, $s(X_1, Y_1), \dots, s(X_{n+1}, Y_{n+1})$ are exchangeable
(if $\hat{\mu}$ is pre-trained on an indep. dataset—i.e., in split conformal prediction)

Conformal prediction ctd.



- ▶ Typical setting: *exchangeable datapoints*.
 - ▶ For a given nonconformity score $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, e.g.,
 $s(x, y) := |y - \hat{\mu}(x)|$, $s(X_1, Y_1), \dots, s(X_{n+1}, Y_{n+1})$ are exchangeable
(if $\hat{\mu}$ is pre-trained on an indep. dataset—i.e., in split conformal prediction)
 - ▶ Hence, the rank of $s(X_{n+1}, Y_{n+1})$ is uniform over $1, \dots, n+1$ (if no ties)

Conformal prediction ctd.

- ▶ Typical setting: *exchangeable datapoints*.
 - ▶ For a given nonconformity score $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, e.g.,
 $s(x, y) := |y - \hat{\mu}(x)|$, $s(X_1, Y_1), \dots, s(X_{n+1}, Y_{n+1})$ are exchangeable
(if $\hat{\mu}$ is pre-trained on an indep. dataset—i.e., in split conformal prediction)
 - ▶ Hence, the rank of $s(X_{n+1}, Y_{n+1})$ is uniform over $1, \dots, n+1$ (if no ties)
 - ▶ So $x \mapsto C(x) = \{y : \text{rank}\{s(x, y) : s_1, \dots, s_n\} \leq \lceil (1 - \alpha)(n + 1) \rceil\}$
satisfies $\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$

Conformal prediction ctd.

- ▶ Typical setting: *exchangeable datapoints*.
 - ▶ For a given nonconformity score $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, e.g.,
 $s(x, y) := |y - \hat{\mu}(x)|$, $s(X_1, Y_1), \dots, s(X_{n+1}, Y_{n+1})$ are exchangeable
(if $\hat{\mu}$ is pre-trained on an indep. dataset—i.e., in split conformal prediction)
 - ▶ Hence, the rank of $s(X_{n+1}, Y_{n+1})$ is uniform over $1, \dots, n+1$ (if no ties)
 - ▶ So $x \mapsto C(x) = \{y : \text{rank}\{s(x, y) : s_1, \dots, s_n\} \leq \lceil (1 - \alpha)(n+1) \rceil\}$
satisfies $\mathbb{P}(Y_{n+1} \in C(X_{n+1})) \geq 1 - \alpha$



Simultaneous Conformal Prediction of Missing Outcomes with Propensity Score ε -Discretization



Yonghoon Lee



Eric Tchetgen Tchetgen


Our problem setting

- Given data

$$(X_1, A_1, Y_1 A_1), \dots, (X_n, A_n, Y_n A_n) \stackrel{\text{iid}}{\sim} P_X \times P_{A|X} \times P_{Y|X},$$


with outcomes *missing at random* (MAR). Thus,

$A_i = 1$: Y_i is observed, $A_i = 0$: Y_i is unobserved.



Our problem setting: Missing At Random

- **Goal:** Construct prediction sets $\{\hat{C}(X_i) : A_i = 0\}$ for the missing outcomes $\{Y_i : A_i = 0\}$



Inferential target

- ▶ With i.i.d./exchangeable data $(X_1, Y_1), \dots, (X_n, Y_n)$ and test input X_{n+1} , standard conformal prediction gives a prediction set $\hat{C}_n(X_{n+1})$ with *marginal coverage*

$$\mathbb{P}(Y_{n+1} \in \hat{C}_n(X_{n+1})) \geq 1 - \alpha.$$

Inferential target

- ▶ With i.i.d./exchangeable data $(X_1, Y_1), \dots, (X_n, Y_n)$ and test input X_{n+1} , standard conformal prediction gives a prediction set $\hat{C}_n(X_{n+1})$ with *marginal coverage*

$$\mathbb{P}(Y_{n+1} \in \hat{C}_n(X_{n+1})) \geq 1 - \alpha.$$

- ▶ **Question:** Under MAR:
 - ▶ What kind of distribution-free inference is possible for missing outcomes?

Inferential target

- ▶ With i.i.d./exchangeable data $(X_1, Y_1), \dots, (X_n, Y_n)$ and test input X_{n+1} , standard conformal prediction gives a prediction set $\hat{C}_n(X_{n+1})$ with *marginal coverage*

$$\mathbb{P}(Y_{n+1} \in \hat{C}_n(X_{n+1})) \geq 1 - \alpha.$$

- ▶ **Question:** Under MAR:
 - ▶ What kind of distribution-free inference is possible for missing outcomes?
 - ▶ Is it possible to go beyond marginal coverage? E.g., have coverage conditional on the test inputs/feature observations with missing outcomes?

Overview of results

- ▶ We consider coverage guarantees of the form

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \widehat{C}(X_i)\} \middle| X_{1:n}, A_{1:n} \right] \geq 1 - \alpha, \quad (1)$$

where $N^{(0)}$ is the number of unobserved labels, and $0/0 := 1$.

- ▶ The **proportion of covered missing outcomes** is on average at least $1 - \alpha$, conditional on $X_{1:n}$ and the missingness pattern $A_{1:n}$.

Overview of results

- ▶ We consider coverage guarantees of the form

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \widehat{C}(X_i)\} \middle| X_{1:n}, A_{1:n} \right] \geq 1 - \alpha, \quad (1)$$

where $N^{(0)}$ is the number of unobserved labels, and $0/0 := 1$.

- ▶ The **proportion of covered missing outcomes** is on average at least $1 - \alpha$, conditional on $X_{1:n}$ and the missingness pattern $A_{1:n}$.
 - ▶ For discrete features X , we construct a procedure that achieves (1).

Overview of results

- ▶ We consider coverage guarantees of the form

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \widehat{C}(X_i)\} \middle| X_{1:n}, A_{1:n} \right] \geq 1 - \alpha, \quad (1)$$

where $N^{(0)}$ is the number of unobserved labels, and $0/0 := 1$.

- ▶ The **proportion of covered missing outcomes** is on average at least $1 - \alpha$, conditional on $X_{1:n}$ and the missingness pattern $A_{1:n}$.
 - ▶ For discrete features X , we construct a procedure that achieves (1).
 - ▶ For general features X , we prove an impossibility result for (1); and then relax it.

Overview of results - continued

- ▶ As a relaxation, we consider

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \widehat{C}(X_i)\} \middle| B_{1:n}, A_{1:n} \right] \geq 1 - \alpha, \quad (2)$$

where $B_i = B_i(X_i)$ is a discretization of X_i (defined soon).

Overview of results - continued

- ▶ As a relaxation, we consider

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \widehat{C}(X_i)\} \mid B_{1:n}, A_{1:n} \right] \geq 1 - \alpha, \quad (2)$$

where $B_i = B_i(X_i)$ is a discretization of X_i (defined soon).

- ▶ **Challenge:** Even though we have MAR ($Y \perp\!\!\!\perp A \mid X$), this does not need to be preserved after discretization (may have $Y \not\perp\!\!\!\perp A \mid B$ for $B = B(X)$).

Overview of results - continued

- ▶ As a relaxation, we consider

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \widehat{C}(X_i)\} \mid B_{1:n}, A_{1:n} \right] \geq 1 - \alpha, \quad (2)$$

where $B_i = B_i(X_i)$ is a discretization of X_i (defined soon).

- ▶ **Challenge:** Even though we have MAR ($Y \perp\!\!\!\perp A \mid X$), this does not need to be preserved after discretization (may have $Y \not\perp\!\!\!\perp A \mid B$ for $B = B(X)$).
- ▶ We introduce a carefully designed **propensity score partitioning scheme**, and show how it can be used to obtain (2)

Table of Contents

Rank-calibration for LLMs

Conformal prediction


Our Methods

Illustrating Our Methods in a Stylized Problem

Empirical illustration


First case: Discrete features

- Discrete features naturally form groups of outcomes $\{Y_i : X_i = x\}$, $x \in \mathcal{X}$.



First case: Discrete features


- ▶ Discrete features naturally form groups of outcomes $\{Y_i : X_i = x\}$, $x \in \mathcal{X}$.



- ▶ Within each group, the outcomes are *exchangeable* conditional on $X_i = x$.


Procedure for discrete features: Naive approach

- Direct method: run split conformal prediction separately for each x .



Procedure for discrete features: Naive approach


- Direct method: run split conformal prediction separately for each x .



- This method attains $\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \hat{C}(X_i)\} \mid X_{1:n}, A_{1:n} \right] \geq 1 - \alpha$.

Procedure for discrete features: Naive approach

- Direct method: run split conformal prediction separately for each x .



- This method attains $\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \hat{C}(X_i)\} \mid X_{1:n}, A_{1:n} \right] \geq 1 - \alpha$.
- However, it can produce infinite-width prediction sets in small groups with $\geq \alpha$ missingness.

Procedure for discrete features: our method

- ▶ Alternative method: simultaneous inference across multiple feature values.

Procedure for discrete features: our method

- ▶ Alternative method: simultaneous inference across multiple feature values.
- ▶ Let
 1. Nonconformity score $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, and $S_i = s(X_i, Y_i)$ if $A_i = 1$
 2. Distinct X values observed: $X'_1, \dots, X'_M$
 3. Indices of datapoints with features equal to X'_k :
 $I_k = \{i \in [n] : X_i = X'_k\}$,
 4. Indices partitioned according to unobserved and observed outcomes,
resp.: $I_k^0 = \{i \in [n] : X_i = X'_k, A_i = 0\}$,
 $I_k^1 = \{i \in [n] : X_i = X'_k, A_i = 1\}$.
 5. Sample sizes $N_k = |I_k|$, $N_k^0 = |I_k^0|$

Procedure for discrete features: our method

- ▶ Alternative method: simultaneous inference across multiple feature values.
- ▶ Let
 1. Nonconformity score $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, and $S_i = s(X_i, Y_i)$ if $A_i = 1$
 2. Distinct X values observed: $X'_1, \dots, X'_M$
 3. Indices of datapoints with features equal to X'_k :
 $I_k = \{i \in [n] : X_i = X'_k\}$,
 4. Indices partitioned according to unobserved and observed outcomes, resp.: $I_k^0 = \{i \in [n] : X_i = X'_k, A_i = 0\}$,
 $I_k^1 = \{i \in [n] : X_i = X'_k, A_i = 1\}$.
 5. Sample sizes $N_k = |I_k|$, $N_k^0 = |I_k^0|$
- ▶ Our prediction set $\hat{C}(x)$:

$$\left\{ y \in \mathcal{Y} : s(x, y) \leq Q_{1-\alpha} \left(\sum_{k=1}^M \sum_{i \in I_k^1} \frac{N_k^0}{N_k N^{(0)}} \delta_{S_i} + \sum_{k=1}^M \frac{(N_k^0)^2}{N_k N^{(0)}} \delta_{+\infty} \right) \right\}. \quad (3)$$

Procedure for discrete features: our method

- ▶ Alternative method: simultaneous inference across multiple feature values.
- ▶ Let
 1. Nonconformity score $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, and $S_i = s(X_i, Y_i)$ if $A_i = 1$
 2. Distinct X values observed: $X'_1, \dots, X'_M$
 3. Indices of datapoints with features equal to X'_k :
 $I_k = \{i \in [n] : X_i = X'_k\}$,
 4. Indices partitioned according to unobserved and observed outcomes, resp.: $I_k^0 = \{i \in [n] : X_i = X'_k, A_i = 0\}$,
 $I_k^1 = \{i \in [n] : X_i = X'_k, A_i = 1\}$.
 5. Sample sizes $N_k = |I_k|$, $N_k^0 = |I_k^0|$
- ▶ Our prediction set $\hat{C}(x)$:

$$\left\{ y \in \mathcal{Y} : s(x, y) \leq Q_{1-\alpha} \left(\sum_{k=1}^M \sum_{i \in I_k^1} \frac{N_k^0}{N_k N^{(0)}} \delta_{S_i} + \sum_{k=1}^M \frac{(N_k^0)^2}{N_k N^{(0)}} \delta_{+\infty} \right) \right\}. \quad (3)$$

- ▶ Idea: conditionally on $X_{1:n}$ and $A_{1:n}$, the distribution of outcomes $Y_{1:n}$ is invariant under the group of permutations that keeps the feature values fixed
- ▶ Provides uniform-width prediction sets for all x values.

Procedure for discrete features: guarantee

Theorem 1

The prediction set (3) satisfies *feature- and missingness-conditional coverage*

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i: A_i=0} I\{Y_i \in \widehat{C}(X_i)\} \mid X_{1:n}, A_{1:n} \right] \geq 1 - \alpha.$$

Discrete features: improvement via partitioning

- ▶ If missingness proportion is high, this can still be conservative.

Discrete features: improvement via partitioning

- ▶ If missingness proportion is high, this can still be conservative.
- ▶ Idea: *Partition* datapoints. For each partition, use $\hat{C}$ on *all datapoints with observed labels* to predict outcomes missing in that partition.

Discrete features: improvement via partitioning

- ▶ If missingness proportion is high, this can still be conservative.
- ▶ Idea: *Partition* datapoints. For each partition, use $\hat{C}$ on *all datapoints with observed labels* to predict outcomes missing in that partition.
- ▶ Since guarantee is feature- and missingness-conditional, this is still valid!

Discrete features: improvement via partitioning

- ▶ If missingness proportion is high, this can still be conservative.
- ▶ Idea: *Partition* datapoints. For each partition, use $\hat{C}$ on *all datapoints with observed labels* to predict outcomes missing in that partition.
- ▶ Since guarantee is feature- and missingness-conditional, this is still valid!
- ▶ Previous methods are at two endpoints: partition is all singletons ("naive method") vs whole set ("our method").

Discrete features: improvement via partitioning


- ▶ If missingness proportion is high, this can still be conservative.
- ▶ Idea: *Partition* datapoints. For each partition, use $\hat{C}$ on *all datapoints with observed labels* to predict outcomes missing in that partition.
- ▶ Since guarantee is feature- and missingness-conditional, this is still valid!
- ▶ Previous methods are at two endpoints: partition is all singletons ("naive method") vs whole set ("our method").
- ▶ Why practically useful? Partition can depend on $X_{1:n}, A_{1:n}$; can aim to ensure small missingness per group.

Procedure for general feature distributions

- ▶ If propensity score $x \mapsto p_{A|X}(x) = \mathbb{P}(A = 1|X = x)$ is known, ε -**discretize** it
- ▶ Let ε be a pre-set discretization level, and $z_k = (1 + \varepsilon)^k / [1 + (1 + \varepsilon)^k]$ for $k \in \mathbb{Z}$

Procedure for general feature distributions

- ▶ If propensity score $x \mapsto p_{A|X}(x) = \mathbb{P}(A = 1|X = x)$ is known, **ε -discretize** it
- ▶ Let ε be a pre-set discretization level, and $z_k = (1 + \varepsilon)^k / [1 + (1 + \varepsilon)^k]$ for $k \in \mathbb{Z}$
- ▶ Partition the feature space into $D_k = \{x : p_{A|X}(x) \in [z_k, z_{k+1})\}$, $\mathcal{B} = \{D_k : k \in \mathbb{Z}\}$.



Pro-CP

- ▶ Show *approximate within-partition exchangeability* of the scores, enabling inference.

Pro-CP

- ▶ Show *approximate within-partition exchangeability* of the scores, enabling inference.
- ▶ **Propensity score discretization-based conformal prediction (pro-CP)**: use on the discretized data $(B_i, A_i, A_i Y_i)_{i \in [n]}$, i.e., $\hat{C}^{\text{pro-CP}}(x)$ is

$$\left\{ y \in \mathcal{Y}, :, s(x, y) \leq Q_{1-\alpha} \left(\sum_{k=1}^M \sum_{i \in I_k^{\mathcal{B},1}} \frac{N_k^{\mathcal{B},0}}{N^{(0)} N_k^{\mathcal{B}}} \cdot \delta_{S_i} + \frac{1}{N^{(0)}} \sum_{k=1}^M \frac{(N_k^{\mathcal{B},0})^2}{N_k^{\mathcal{B}}} \cdot \delta_{+\infty} \right) \right\}.$$

Pro-CP

- ▶ Show *approximate within-partition exchangeability* of the scores, enabling inference.
- ▶ **Propensity score discretization-based conformal prediction (pro-CP)**: use on the discretized data $(B_i, A_i, A_i Y_i)_{i \in [n]}$, i.e., $\hat{C}^{\text{pro-CP}}(x)$ is

$$\left\{ y \in \mathcal{Y}, :, s(x, y) \leq Q_{1-\alpha} \left(\sum_{k=1}^M \sum_{i \in I_k^{\mathcal{B},1}} \frac{N_k^{\mathcal{B},0}}{N^{(0)} N_k^{\mathcal{B}}} \cdot \delta_{S_i} + \frac{1}{N^{(0)}} \sum_{k=1}^M \frac{(N_k^{\mathcal{B},0})^2}{N_k^{\mathcal{B}}} \cdot \delta_{+\infty} \right) \right\}.$$

Theorem 2

If $0 < p_{A|X}(X) < 1$ a.s., then $\hat{C}^{\text{pro-CP}}$ has *propensity score discretized feature- and missingness-conditional coverage*:

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i: A_i=0} I\{Y_i \in \hat{C}^{\text{pro-CP}}(X_i)\} \mid B_{1:n}, A_{1:n} \right] \geq 1 - \alpha - \varepsilon.$$

Pro-CP

- ▶ Show *approximate within-partition exchangeability* of the scores, enabling inference.
- ▶ **Propensity score discretization-based conformal prediction (pro-CP)**: use on the discretized data $(B_i, A_i, A_i Y_i)_{i \in [n]}$, i.e., $\hat{C}^{\text{pro-CP}}(x)$ is

$$\left\{ y \in \mathcal{Y}, :, s(x, y) \leq Q_{1-\alpha} \left(\sum_{k=1}^M \sum_{i \in I_k^{\mathcal{B},1}} \frac{N_k^{\mathcal{B},0}}{N^{(0)} N_k^{\mathcal{B}}} \cdot \delta_{S_i} + \frac{1}{N^{(0)}} \sum_{k=1}^M \frac{(N_k^{\mathcal{B},0})^2}{N_k^{\mathcal{B}}} \cdot \delta_{+\infty} \right) \right\}.$$

Theorem 2

If $0 < p_{A|X}(X) < 1$ a.s., then $\hat{C}^{\text{pro-CP}}$ has *propensity score discretized feature- and missingness-conditional coverage*:

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i: A_i=0} I\{Y_i \in \hat{C}^{\text{pro-CP}}(X_i)\} \mid B_{1:n}, A_{1:n} \right] \geq 1 - \alpha - \varepsilon.$$

- ▶ The error from discretization is $\leq \varepsilon$, for *any* n and $\#$ of missing outcomes.

Pro-CP with estimated propensity score

- ▶ If the propensity score is unknown, run pro-CP with an estimator $\hat{p}_{A|X}$ of $p_{A|X}$.

Pro-CP with estimated propensity score

- ▶ If the propensity score is unknown, run pro-CP with an estimator $\hat{p}_{A|X}$ of $p_{A|X}$.

Theorem 3

Suppose $0 < p_{A|X}(X) < 1$ and $0 < \hat{p}_{A|X}(X) < 1$ almost surely. Then pro-CP run with $\hat{p}_{A|X}$ satisfies

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \hat{C}^{\text{pro-CP}}(X_i)\} \mid B_{1:n}, A_{1:n} \right] \geq 1 - \alpha - (\varepsilon + \delta_{\hat{p}_{A|X}} + \varepsilon \delta_{\hat{p}_{A|X}}),$$

where

$$\delta_{\hat{p}_{A|X}} = e^{2\|\log f_{p,\hat{p}}\|_\infty} - 1, \quad f_{p,\hat{p}}(x) = \frac{p_{A|X}(x)/(1-p_{A|X}(x))}{\hat{p}_{A|X}(x)/(1-\hat{p}_{A|X}(x))}.$$

Pro-CP with estimated propensity score

- ▶ If the propensity score is unknown, run pro-CP with an estimator $\hat{p}_{A|X}$ of $p_{A|X}$.

Theorem 3

Suppose $0 < p_{A|X}(X) < 1$ and $0 < \hat{p}_{A|X}(X) < 1$ almost surely. Then pro-CP run with $\hat{p}_{A|X}$ satisfies

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i:A_i=0} I\{Y_i \in \hat{C}^{\text{pro-CP}}(X_i)\} \mid B_{1:n}, A_{1:n} \right] \geq 1 - \alpha - (\varepsilon + \delta_{\hat{p}_{A|X}} + \varepsilon \delta_{\hat{p}_{A|X}}),$$

where

$$\delta_{\hat{p}_{A|X}} = e^{2\|\log f_{p,\hat{p}}\|_\infty} - 1, \quad f_{p,\hat{p}}(x) = \frac{p_{A|X}(x)/(1-p_{A|X}(x))}{\hat{p}_{A|X}(x)/(1-\hat{p}_{A|X}(x))}.$$

- ▶ The error from estimation does not grow with the number of missing outcomes.

New result underlying pro-CP guarantee

- ▶ Balancing property of the propensity score [Rosenbaum and Rubin (1983)]: the missingness is independent of the outcome conditional on the propensity: $A \perp\!\!\!\perp Y \mid p_{A|X}$.

New result underlying pro-CP guarantee

- ▶ Balancing property of the propensity score [Rosenbaum and Rubin (1983)]: the missingness is independent of the outcome conditional on the propensity: $A \perp\!\!\!\perp Y \mid p_{A|X}$.
- ▶ We show *approximate version*: dist. of $s(X, Y)$ close for $A = 0, 1$ given small range of $p_{A|X}$

Lemma (Bounded prop. score $\Rightarrow$ cond. dist. of obs. and missing are close)

Suppose that $(X, Y, A) \sim P_X \times P_{Y|X} \times \text{Bernoulli}(p_{A|X})$ on $\mathcal{X} \times \mathcal{Y} \times \{0, 1\}$, and that for a set $B \subset \mathcal{X}$ and $t \in (0, 1)$, $\varepsilon \geq 0$,

$$t \leq \frac{p_{A|X}(x)}{1 - p_{A|X}(x)} \leq t(1 + \varepsilon), \text{ for all } x \in B.$$

Let $s : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$ be any measurable function and let $S = s(X, Y)$.

Then
$$d_{TV}(P_{S|A=1, X \in B}, P_{S|A=0, X \in B}) \leq \varepsilon.$$

Table of Contents

Rank-calibration for LLMs

Conformal prediction

Our Methods

Illustrating Our Methods in a Stylized Problem

Empirical illustration

Application to simultaneous inference on ITEs

- Consider a potential outcomes model

$$(X_i, T_i, Y_i(0), Y_i(1))_{1 \leq i \leq n} \stackrel{\text{iid}}{\sim} P_X \times P_{T|X} \times P_{Y(1)|X} \times P_{Y(0)|X},$$

where we observe $(X_i, T_i, T_i Y_i(1) + (1 - T_i) Y_i(0))_{1 \leq i \leq n}$.

Application to simultaneous inference on ITEs

- Consider a potential outcomes model

$$(X_i, T_i, Y_i(0), Y_i(1))_{1 \leq i \leq n} \stackrel{\text{iid}}{\sim} P_X \times P_{T|X} \times P_{Y(1)|X} \times P_{Y(0)|X},$$

where we observe $(X_i, T_i, T_i Y_i(1) + (1 - T_i) Y_i(0))_{1 \leq i \leq n}$.

- Applying pro-CP, we can construct $\hat{C}^{\text{counterfactual}}$ such that

$$\mathbb{E} \left[\frac{1}{N(0)} \sum_{i: T_i=0} \mathbf{1} \left\{ Y_i(1) \in \hat{C}^{\text{counterfactual}}(X_i) \right\} \mid B_{1:n}, T_{1:n} \right] \geq 1 - \alpha$$

Application to simultaneous inference on ITEs

- ▶ Consider a potential outcomes model

$$(X_i, T_i, Y_i(0), Y_i(1))_{1 \leq i \leq n} \stackrel{\text{iid}}{\sim} P_X \times P_{T|X} \times P_{Y(1)|X} \times P_{Y(0)|X},$$

where we observe $(X_i, T_i, T_i Y_i(1) + (1 - T_i) Y_i(0))_{1 \leq i \leq n}$.

- ▶ Applying pro-CP, we can construct $\hat{C}^{\text{counterfactual}}$ such that

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i: T_i=0} \mathbf{1} \left\{ Y_i(1) \in \hat{C}^{\text{counterfactual}}(X_i) \right\} \mid B_{1:n}, T_{1:n} \right] \geq 1 - \alpha$$

- ▶ By letting $\hat{C}_i^{\text{ITE}} = \{y - Y_i(0) : y \in \hat{C}^{\text{counterfactual}}(X_i)\}$, we obtain prediction sets for individual treatment effects

$$\mathbb{E} \left[\frac{1}{N^{(0)}} \sum_{i \in I_{T=0}} \mathbf{1} \left\{ (Y_i(1) - Y_i(0)) \in \hat{C}_i^{\text{ITE}} \right\} \mid B_{1:n}, T_{1:n} \right] \geq 1 - \alpha$$

Table of Contents

Rank-calibration for LLMs

Conformal prediction

Our Methods

Illustrating Our Methods in a Stylized Problem

Empirical illustration

Simulation 1

Weighted conformal (Tibshirani et al., 2019) vs pro-CP: marginal vs conditional coverage

Simulation 1

Weighted conformal (Tibshirani et al., 2019) vs pro-CP: marginal vs conditional coverage

1. $X \sim \text{Unif}[0, 10]$, $Y \mid X \sim N(X, (3 + X)^2)$, $A \mid X \sim \text{Bernoulli}(p_{A|X}(X))$

Simulation 1

Weighted conformal (Tibshirani et al., 2019) vs pro-CP: marginal vs conditional coverage

1. $X \sim \text{Unif}[0, 10]$, $Y \mid X \sim N(X, (3 + X)^2)$, $A \mid X \sim \text{Bernoulli}(p_{A|X}(X))$
2. (1) : $p_{A|X}(x) = 0.9 - 0.02x$, (2) : $p_{A|X}(x) = 0.8 - 0.1(1 + 0.1x) \sin 3x$

Simulation 1

Weighted conformal (Tibshirani et al., 2019) vs pro-CP: marginal vs conditional coverage

1. $X \sim \text{Unif}[0, 10]$, $Y \mid X \sim N(X, (3 + X)^2)$, $A \mid X \sim \text{Bernoulli}(p_{A|X}(X))$
2. (1) : $p_{A|X}(x) = 0.9 - 0.02x$, (2) : $p_{A|X}(x) = 0.8 - 0.1(1 + 0.1x) \sin 3x$
3. Fit OLS with $n_{\text{train}} = 500$, $s(x, y) = |y - \hat{\mu}(x)|$

Simulation 1

Weighted conformal (Tibshirani et al., 2019) vs pro-CP: marginal vs conditional coverage

1. $X \sim \text{Unif}[0, 10]$, $Y \mid X \sim N(X, (3 + X)^2)$, $A \mid X \sim \text{Bernoulli}(p_{A|X}(X))$
2. (1) : $p_{A|X}(x) = 0.9 - 0.02x$, (2) : $p_{A|X}(x) = 0.8 - 0.1(1 + 0.1x) \sin 3x$
3. Fit OLS with $n_{\text{train}} = 500$, $s(x, y) = |y - \hat{\mu}(x)|$
4. 500 trials, $n = 500$, Pro-CP $\varepsilon = 0.1$, $\alpha = 0.2$, partition of size 10;

Simulation 1

Weighted conformal (Tibshirani et al., 2019) vs pro-CP: marginal vs conditional coverage

1. $X \sim \text{Unif}[0, 10]$, $Y \mid X \sim N(X, (3 + X)^2)$, $A \mid X \sim \text{Bernoulli}(p_{A|X}(X))$
2. (1) : $p_{A|X}(x) = 0.9 - 0.02x$, (2) : $p_{A|X}(x) = 0.8 - 0.1(1 + 0.1x) \sin 3x$
3. Fit OLS with $n_{\text{train}} = 500$, $s(x, y) = |y - \hat{\mu}(x)|$
4. 500 trials, $n = 500$, Pro-CP $\varepsilon = 0.1$, $\alpha = 0.2$, partition of size 10;
5. Given $X_{1:n}, A_{1:n}$, 100x gen $(X'_i, Y'_i)_{1 \leq i \leq n} \mid B_i \sim P_{X|B} \times P_{Y|X}$, $n = 500$

Simulation 1

Weighted conformal (Tibshirani et al., 2019) vs pro-CP: marginal vs conditional coverage

1. $X \sim \text{Unif}[0, 10]$, $Y | X \sim N(X, (3 + X)^2)$, $A | X \sim \text{Bernoulli}(p_{A|X}(X))$
2. (1) : $p_{A|X}(x) = 0.9 - 0.02x$, (2) : $p_{A|X}(x) = 0.8 - 0.1(1 + 0.1x) \sin 3x$
3. Fit OLS with $n_{\text{train}} = 500$, $s(x, y) = |y - \hat{\mu}(x)|$
4. 500 trials, $n = 500$, Pro-CP $\varepsilon = 0.1$, $\alpha = 0.2$, partition of size 10;
5. Given $X_{1:n}, A_{1:n}$, 100x gen $(X'_i, Y'_i)_{1 \leq i \leq n} | B_i \sim P_{X|B} \times P_{Y|X}$, $n = 500$




Illustration on diabetes dataset (Efron et al., 2004)

- ▶ X : ten features (age, bmi, LDL/HDL, ...) of patients (sample sizes: train: 142; calibration+test: 300)
- ▶ A : missingness generated from a known logistic model
- ▶ Y : a measure of disease progression one year after baseline




Illustration on diabetes dataset (Efron et al., 2004): II





Illustration on JOBS II dataset (Imai et al., 2010)

- ▶ X : job seekers: $n_{\text{train}} = 379$, $n = 500$; with 14 demographic features
- ▶ A : job skills workshop (to evaluate our methods, simulate via logistic model; estimate via RF)
- ▶ $Y(0)$, $Y(1)$: pre- and post-treatment depression measure



Summary for conformal prediction

- ▶ Introduced Pro-CP, a method for simultaneous prediction of multiple missing outcomes, and provided coverage guarantees
- ▶ Pro-CP2: stronger squared error miscoverage error control

Thank you!



SIM NS
FOUNDATION


ALFRED P. SLOAN
FOUNDATION



Jailbreaking Black Box Large Language Models in Twenty Queries



Patrick Chao, Alexander Robey,
Edgar Dobriban, Hamed Hassani, George J. Pappas, Eric Wong
University of Pennsylvania



JailbreakBench: An Open Robustness Benchmark for Jailbreaking Large Language Models



Patrick Chao^{*1}, Edoardo DeBenedetti^{*2}, Alexander Robey^{*1}, Maksym Andriushchenko^{*3},
Francesco Croce³, Vikash Sehwal⁴, Edgar Dobriban¹, Nicolas Flammarion³,
George J. Pappas¹, Florian Tramèr², Hamed Hassani¹, Eric Wong¹

[Leaderboards](#)[Paper](#)[Contribute](#)[Library](#)[Behaviors](#)[Jailbreak artifacts](#)

JAILBREAKBENCH

Other work (ctd)



One-Shot Safety Alignment for Large Language Models via Optimal Dualization

Xinmeng Huang*
Osbert Bastani

Shuo Li*
Hamed Hassani

Edgar Dobriban
Dongsheng Ding†

Watermarking Language Models with Error Correcting Codes

Patrick Chao, Edgar Dobriban, Hamed Hassani*



Evaluating the Performance of Large Language Models via Debates

Behrad Moniri

Hamed Hassani

Edgar Dobriban